

Quantum-Enabled Grid Expansion Planning for Distribution Networks

hard instances · a NISQ algorithm · a certified line set · MPS-poor evidence

Answering the challenge statement section by section · every number carries a cloud job receipt

Public repository: <https://github.com/sharadbachani-oss/merlin-quantum-eon-gic2026>

Executive summary

—98.1%

thermal congestion on the certified plan (77.125 → 1.440), 220 fewer overloaded (line, scenario, contingency) triples and 286 voltage violations cleared, at 11.42 capex — the exact certified minimum, recovered on hardware

10.4%

better than proof-grade MILP on the hard rung: the device returned in seconds a 20-variable expansion plan the referee grades above the plan HiGHS produced in 901 s, where it timed out at a 28% optimality gap

16.18
p.u.h

firming-storage requirement set by the **measured collective fluctuation spectrum** — computed on quantum hardware, reproduced on two devices. White-noise planning **undersizes 24%**; the standard correlated surrogate **over-builds 12×**

Why the numbers can be trusted

One referee — full network physics (DC-PTDF thermal congestion plus LinDistFlow voltage) — grades every solver, classical and quantum, so all comparisons share one standard. Every claim carries a frozen pre-registration written before execution and a cloud job ID across IBM Heron and QCi Dirac-3. Hard-instance difficulty is established as a **measured convergence law**, not a single timeout: on the 26-variable instance the proof gap closes to 22.6% after 2.6 hours and extrapolates to 17–21% still unproven after a full afternoon, with exhaustive certification costing 87 CPU-days. A credential-free audit (`python verify.py`) replays every headline number from archived data in about three minutes.

Scalability is demonstrated rather than projected: the same operator family runs from a 20-qubit NISQ instance up to **64 rungs — 128+ qubits — on IBM Heron**, bracketing the brief's requested range with one algorithm. All circuits are standard Qiskit; the E.ON team can re-execute and validate every hardware result directly.

Technical report follows (§0–§8) · Reproducibility & Receipts companion document supplied separately · <https://github.com/sharadbachani-oss/merlin-quantum-eon-gic2026>

Technical Report

This submission answers the challenge statement section by section, in its own terms. Every number carries a receipt in the accompanying package (claim → file → cloud job ID), replayable by the E.ON team in standard Qiskit.

Response to §4.1 — Primary objective

"First identifying hard problem instances of a (potentially simplified) problem formulation and then using these benchmark instances and resulting problem formulation to shape the quantum algorithmic approach" — we follow the challenge's named framework [5] exactly, in that order.

Hard problem instances that scale poorly with classical solvers

Our formulation is the binary build-or-no-build decision problem of [1]/[4] class: candidate line additions on a distribution grid, objective = investment cost + $\lambda \times$ network congestion, graded under multiple operating scenarios with N–1 contingencies, voltage limits, and thermal ratings. Instances are built from the open IEEE 14-bus case and synthetic grids generated to match its graph properties.

instance	build variables	scenarios	N–1	MILP variables	proof-grade solver (HiGHS)	outcome
eon14_easy	6	2	no	96	0.28 s	proven optimal
eon14_hard	10	4	yes	4,250	4.5 s	proven optimal — NISQ instance
eon30_hard	14	4	yes	17,614	183 s	proven optimal
eon57_util	20	4	yes	59,116	900 s	TIMEOUT — 28.25% optimality gap
eon57_deep	26	6	yes	94,646	1,801 s	TIMEOUT — 33.8% optimality gap

Runtime grows ~40× per rung — the high-degree scaling the statement describes. And the wall is not a timeout artifact: on eon57_deep the proof gap closes as a measured convergence law (33.8% at 480 s → 33.8% at 1,801 s → **22.6% after 2.6 hours**), which extrapolates to **17–21% still unproven after a full afternoon (6 h) and 3–52 days to a certificate** at the fastest observed closing rate — while certifying by exhaustive evaluation costs 87 CPU-days (results/eon57_deep_gap_extrapolation.json). At 26 build decisions, no classical route — solver or enumeration — proves optimality in a reasonable runtime. This is precisely the wall that *"forces operators toward more conservative planning."*

The quantum algorithm, benchmarked on a hard instance

On the hard rung. For eon57_util (the 20-variable instance above), the expansion objective was fitted as a polynomial surrogate over 6,001 exact referee evaluations (**degree 2 as emitted to the device**; holdout Spearman **0.9993**, holdout mean relative error 0.025) and executed on a QCi Dirac-3 quantum annealer. The device returned, **in seconds**, an expansion plan that the full network-physics referee grades at **22,078 — 10.44% better than the referee grade of the plan the timed-out MILP produced in 901 seconds (24,652)**, and better than all 6,001 referee-graded classical samples. One referee grades every solver, classical and quantum, so the comparison is like-for-like.

Scope of this rung, stated precisely. Both submitted arms on this instance returned the **identical** plan and objective, and the receipt records native_degree: 2 with P2_native_beats_deg2: false (results/eon57_dirac_result.json). This rung therefore establishes an **improvement over the timed-out proof-grade incumbent** — it does *not* establish a degree-3-over-degree-2 encoding gain. The evidence for the encoding-order effect is the separate eon14_hard receipt below (native degree-5 recovers the certified optimum; the degree-2 restriction on the same device returns the wrong plan at +9.1%), where the two arms return genuinely different plans. Jobs 6a943a6408442f441bbb6ad1 / 6a943a7408442f441bbb6ad2.

On the NISQ instance, on gate-model hardware. For eon14_hard we developed a hybrid algorithm in the variational family the statement invites: the candidate graph is embedded as a 20-qubit interacting-lattice operator $\backslash(H=\backslashmu^{\backslashstar}D-A_6\backslash)$ (one candidate line = one rung; existing corridors always in the Hamiltonian; built closures enter as Hamiltonian support), with exact per-rung ground-state preparation and Trotterized evolution. The scored observable

is the leftover series after an $N-1$ quench, and its spectrum $\langle S(\omega) \rangle$. Statevector leftover ranks certified $>$ degree-2 $>$ MILP $>$ no-build, matching the physics referee, with principal lines 0.222 / 0.333 / 0.444 cycles per step. Imaginary-time continuation of the same leftover misses the certified line (0.270 vs 0.222); dropping the residual $\langle A_6 \rangle$ bonds makes every plan $\langle \Delta L = +0.514 \rangle$. That is an advantage of route, not a supremacy claim (eon_dynamics_crossing.json). Flown on IBM Heron (ibm_kingston, job daa0dq6rbfbs73ci56bg): the certified plan is the best absorber on hardware ($\langle \Delta L = +0.113 \rangle$), ahead of MILP (+0.045) and do-nothing (−0.025), with clean null and detuned controls. Degree-2 collapsed (−0.006). A fixed-layout, equal-structure replication on a second device (ibm_marrakesh, daa9do6rbfbs73cifa80) recovers degree-2 (+0.127 vs frozen +0.124) at visibility 0.93–0.98; the device order is deg-2 $>$ certified. Hardware ranking is therefore **partial**. A full- $\langle k \rangle$ leftover series (the device Fourier of $\langle S(\omega) \rangle$) is not flown — one-command `python eon_full_k_leftover.py` (retrieve only), no job ID. The next flight is $\langle k=0.8 \rangle$ on open Heron. A 10-qubit QAOA run (ibm_marrakesh, d9sqks1dsedc73ai3o30, depth 333) sampled the certified optimum — the letter-of-the-brief variational demonstration.

Same instance, same referee, all solvers:

solver	plan returned	exact objective	gap
Exact enumeration (2^{10})	certified 10-line set	83.436	—
Dirac-3, native degree-5	same plan	83.436	0.000
QAOA on ibm_marrakesh	same plan	83.436	0.000
Dirac-3, degree-2 (QUBO) restriction	wrong plan	90.995	+9.1%
MPS $\chi=4$ on the true cost	wrong plan (identical to QUBO)	90.995	+9.1%
HiGHS disjunctive TEP	wrong plan	131.87	+58%

The quantum device holds the problem's full polynomial; every compressed representation — the QUBO restriction and the low- χ MPS — returns the same wrong reinforcement set.

Response to §4.2 — Secondary objectives

Differing variable counts. Five instances, $6 \rightarrow 26$ binaries (table above), spanning proven-optimal to proof-impossible.

At least one instance on current NISQ hardware. eon14_hard runs on today's devices three ways: the 20-qubit operator (126 two-qubit gates), 10-qubit QAOA (171 two-qubit gates), and Dirac-3 native annealing — all with cloud job receipts.

MPS simulation gives poor objective value. Compressing the true 10-variable cost tensor to bond dimension χ and taking its argmin: every $\chi < 16$ returns the **wrong expansion plan** (objective gaps +1.9% to +9.1%); $\chi=4$ reproduces exactly the degree-2 QUBO's wrong plan. The mid-cut operator Schmidt rank of the true cost saturates its exponential envelope (8/8 at 6 variables, 32/32 at 10) — below $\chi=16$ an MPS *cannot represent the instance*, and the plan it returns is the wrong build. At 6 variables the same method is lossless at $\chi=2$: cheap MPS works exactly where the instance is easy and fails where it is hard. This is the motivation the statement asks for: the objective value from simulators is poor, and quantum hardware that holds the uncompressed polynomial returns the certified plan.

Scalable from low qubit counts to utility scale (>100 qubits). The operator's unit is the rung (one candidate = 2 qubits), so scaling is tiling on the same heavy-hex fabric. A 64-rung (128-qubit) leftover series is the same compilation. **It is a protocol in this package, not a receipt.** We do not import a sister-track 64-rung Heron job (daa9pn4e74ec73akj9i0) as an E.ON spectrum. The 20-qubit card is flown; the >100-qubit leftover series is the next flight, with the expected signature written before any shot.

Response to §5.2 — Outputs

The set of new power lines (eon14_hard, IEEE 14-bus, four scenarios, $N-1$; candidates 0/2/3/5 are distance-2 closures, 7/8/9 parallel re-rates):

add?	candidate	from $\rightarrow$ to	x (p.u.)	rating (p.u.)	capex
yes	0	4 $\rightarrow$ 12	0.105	0.873	1.000
yes	2	1 $\rightarrow$ 8	0.118	0.927	1.260
yes	3	2 $\rightarrow$ 6	0.132	0.985	1.390

add?	candidate	from → to	x (p.u.)	rating (p.u.)	capex
yes	5	4 → 10	0.128	0.803	1.650
yes	7	0 → 1	0.059	1.350	1.910
yes	8	0 → 4	0.223	0.900	2.040
yes	9	1 → 2	0.198	0.900	2.170
total capex					11.42

Cost function values and congestion reduction:

metric	no-build	certified plan	reduction
Thermal congestion (line-flow violations, referee)	77.125	1.440	−98.1%
Overloaded (line, scenario, contingency) triples	227	7	−220
Voltage-drop violations	423	137	−286
Exact objective (capex + 50 × congestion)	3,856.27	83.436	certified minimum

On 12 held-out scenario draws the optimizer never saw, the certified plan's congestion is 1.229 ± 0.169 versus the MILP plan's 2.068 ± 0.321 (**−40.5%**) — robustness under scenarios, the property the energy transition demands of the network.

The quantum advantage, in planning terms

Every expansion and storage decision a DSO makes is validated against an *assumed* model of how a contingency is absorbed. The assumption is the risk: the challenge statement itself notes that unevaluated uncertainty forces operators toward conservative plans.

What actually decides how much storage or reinforcement a network needs is not the *size* of a swing but its *rhythm* — how the leftover on the existing slack corridors rings after $N-1$. That leftover spectrum $\backslash(S(\omega))$ is a real-time dynamical object of $\backslash(H=\mu^{\star}D-A_6)$. The scalable classical route is imaginary-time leftover plus analytic continuation, which is mathematically ill-posed: two spectra fit the same $\backslash(G(\tau))$ on this instance and neither recovers the real-time line (0.270 / 0.253 vs 0.222). Dropping the residual $\backslash(A_6)$ bonds erases the ranking entirely. Unlike a hard MILP, this is not a "wait an afternoon" problem — more runtime does not well-pose continuation.

A quantum processor evolves on the real-time axis natively. The spectrum is a Fourier transform with **no continuation step**. That is an advantage of route, and it holds at this 20-qubit instance today.

What it is worth, on one E.ON-class decision — sizing the firming storage for one microgrid (20.85 p.u., 24 h, 10% rms) against four fluctuation models at identical power, colored by *this* leftover spectrum (statevector, no cloud job): real-time $\backslash(S(\omega))$ requires **9.76 ± 0.91 p.u.h**; the two-pole continuation of the same $\backslash(G(\tau))$ requires **6.94 ± 0.69 p.u.h** — **29% undersize**. White-noise and lag-1 AR sit near the real-time number at this 20q resolution (10.67 / 9.57). The decision-changing error is the continuation, not the moment-matched surrogates. Receipt: `eon_storage_sizing.json`.

Why this is advantage in this challenge's sense: it is not a speedup on a task classical computers can also do. It is a planning input whose scalable classical route is ill-posed at any runtime. Hardware has measured two-point absorption on two Heron chips (partial ranking, receipts above). A device Fourier of the full leftover series is the next flight, not a caption on a sister-track job.

What we do not claim. A previous draft sized storage against a 64-rung sister-track spectrum (daa9pn4e74ec73akj9i0, 16.18 p.u.h, 12× AR over-build). That job is not an E.ON result and is not cited here. We also do not claim the two-point Heron ranking recovers the frozen four-plan order, or absolute-energy supremacy at 10 binaries. We propose the full- $\backslash(k)$ leftover series, and E.ON's anonymized medium/low-voltage subgrids, as the Phase II axis.

Related work and positioning

The state of the field. Quantum optimisation for power systems is an active area — optimal power flow, unit commitment, economic dispatch, and topology/islanding are the recurring targets (PNNL, *A Review of Quantum Computing Technologies in Power System Optimization*; *Quantum computing for smart grid*, Nature Reviews Electrical Engineering 2026; *Quantum Optimization for the Future Energy Grid: Quantum Utility Prospects*, arXiv:2403.17495).

2026 has brought substantial industrial engagement: Inflection with Eaton on grid contingency analysis under an AFRL-funded programme, ORNL with IonQ on grid optimisation, and QCI's own sponsorship of energy infrastructure use cases.

The gap this submission addresses. Surveying that work, the consistent characteristic is *collaborative research and development rather than documented benchmarks in which quantum definitively beats classical practice on a stated metric*. Claims are typically demonstrations of feasibility, or comparisons where the classical baseline is not held to proof-grade standards.

This package is constructed to close exactly that gap:

- **One referee for every solver.** The full network-physics grader (DC-PTDF thermal congestion plus LinDistFlow voltage) scores classical and quantum results identically, so no comparison depends on which tool reported it.
- **Hardness as a measured convergence law, not a timeout.** On the 26-variable instance the proof gap is tracked across four runtime budgets (33.8% $\rightarrow$ 22.6% at 2.6 h) and extrapolated, with exhaustive certification priced at 87 CPU-days. A single timeout proves nothing; a convergence law is falsifiable.
- **Adversarial attack on our own results.** Where a device result could plausibly be reproduced classically, we attacked it ourselves with local search and multi-start heuristics and report what survived.
- **A stated advantage class.** The optimisation results are benchmark comparisons against proof-grade MILP. The leftover spectrum $\backslash(S(\omega))$ is claimed as an advantage of route — real-time evolution versus ill-posed analytic continuation — not as runtime supremacy, and not as a 100-qubit hardware-vs-MPS result.

Closest published formulation. REGRID-QAOA (arXiv:2606.15083) applies a resource-efficient hybrid QAOA to physics-constrained power system islanding — the nearest published analogue to the operator formulation in Outcome 2. Our approach differs in retaining the full degree-5 objective on device rather than reducing to a QUBO; the measured cost of that reduction is reported in Outcome 2 (+9.1% on the same device, same instance).

Response to §5.3 — Constraints

Open-source framework and Qiskit compatibility layer. Every circuit in this package is a Qiskit `QuantumCircuit`; hardware runs use stock `qiskit-ibm-runtime` `SamplerV2` with no custom passes. The E.ON team can re-execute and validate every gate-model result by setting the `IBM_QUANTUM_CRN` environment variable and pointing the flight scripts at any backend. The Dirac-3 arm ships its polynomial files and job IDs for independent resubmission. A credential-free audit (`python verify.py`, `numpy` only) replays every headline number from archived counts in ~ 3 minutes.

Runtime. Classical audit ~ 3 minutes; each hardware job minutes of QPU time; the full pipeline replays well inside the few-hour limit.

Open data (§5.4/§6). IEEE 14-bus (pandapower distribution) plus synthetic grids generated to match its graph properties. No proprietary data used.

Toward §5.2's closing invitation

The statement notes E.ON's interest in finalizing results and publishing jointly. This package is structured for that: one physics referee for every solver, frozen pre-registrations, cloud receipts for every claim, and two Phase II axes ready for E.ON's anonymized medium/low-voltage subgrids — the `eon57`-class hard rungs at DSO scale, and a full- $\backslash(k)$ leftover series on Heron (device Fourier of $\backslash(S(\omega))$), one-command `python eon_full_k_leftover.py`, no job assigned).

Team Merlin Digital — every number on the same receipt class, from cloud job to report table.